

Projects

Johannes Ibald

November 07, 2020

Contents

1	Resource Latex	1
1.1	A Section	2
1.2	Cool Pictures	3
2		6
3	•	7
4	•	8
5	•	9
6	Creating Realistic Rendering Effects	10
6.1	Understanding graphics shaders	10
6.1.1	osg::Shader	11
6.2	Using uniforms	12
6.2.1	osg::Uniform class	12
6.2.2	Time for Action page 154	13
6.2.3	What just happened?	13
6.3	Working with the geometry shader	14
6.3.1	Time for action - Generating a Bezier curve	14
6.3.2	What just happened?	14
7	Viewing the World	16
7.0.1	From world to screen	16
7.1	The Camera class	18
7.1.1	Rendering order of cameras	20
7.1.2	Time for action creating an HUD camera	21
7.1.3	WTF just happened?	21
7.1.4	using a single viewer	21
7.1.5	Digging into the simulation loop	22
7.1.6	Time for action - custom simulation loop	23

7.1.7	What on earth just happended?	23
7.1.8	Using a composite viewer	24
7.1.9	Time for action - rendering more scenes at one time . .	24
7.1.10	What the fuck just happened?	25
7.1.11	Changing global display settings	25
7.1.12	Time for action - enabling global multisampling	26
7.1.13	wth just happened?	26
7.1.14	Stereo visualization	26
7.1.15	Time for action - rendering naglyph stereo scenes . . .	27
7.1.16	wtf just happened?	27
7.1.17	Rendering to textures	27
7.1.18	Frame buffer, pixel buffer, and FBO	28
7.1.19	TFA - drawing aircrafts on a loaded terrain	29

Chapter 1

Resource Latex

Hello, here is some text without a meaning. This text should show what a printed text will look like at this place. $\sin^2(\alpha) + \cos^2(\beta) = 1$. If you read this text, you will get no information $E = mc^2$. Really? Is there no information? Is there a difference between this text and some nonsense like “Huardest gefburn”? Kjift – not at all! A blind text like this gives you information about the selected font, how the letters are written and an impression of the look. $\sqrt[n]{a} \cdot \sqrt[n]{b} = \sqrt[n]{ab}$. This text should contain all letters of the alphabet and it should be written in of the original language. $\frac{\sqrt[n]{a}}{\sqrt[n]{b}} = \sqrt[n]{\frac{a}{b}}$. There is no need for special content, but the length of words should match the language. $a\sqrt[n]{b} = \sqrt[n]{a^n b}$. Hello, here is some text without a meaning. $d\Omega = \sin\vartheta d\vartheta d\varphi$. This text should show what a printed text will look like at this place. If you read this text, you will get no information. Really? Is there no information? Is there a difference between this text and some nonsense like “Huardest gefburn”? Kjift – not at all! A blind text like this gives you information about the selected font, how the letters are written and an impression of the look. $\sin^2(\alpha) + \cos^2(\beta) = 1$. This text should contain all letters of the alphabet and it should be written in of the original language $E = mc^2$. There is no need for special content, but the length of words should match the language. $\sqrt[n]{a} \cdot \sqrt[n]{b} = \sqrt[n]{ab}$. Hello, here is some text without a meaning. $\frac{\sqrt[n]{a}}{\sqrt[n]{b}} = \sqrt[n]{\frac{a}{b}}$. This text should show what a printed text will look like at this place. $a\sqrt[n]{b} = \sqrt[n]{a^n b}$. If you read this text, you will get no information. $d\Omega = \sin\vartheta d\vartheta d\varphi$. Really? Is there no information? Is there a difference between this text and some nonsense like “Huardest gefburn”? Kjift – not at all! A blind text like this gives you information about the selected font, how the letters are written and an impression of the look. This text should contain all letters of the alphabet and it should be written in of the original language. There is no need for special content, but the length of words should

match the language. $\sin^2(\alpha) + \cos^2(\beta) = 1$. Hello, here is some text without a meaning $E = mc^2$. This text should show what a printed text will look like at this place. $\sqrt[n]{a} \cdot \sqrt[n]{b} = \sqrt[n]{ab}$. If you read this text, you will get no information. $\frac{\sqrt[n]{a}}{\sqrt[n]{b}} = \sqrt[n]{\frac{a}{b}}$. Really? Is there no information? Is there a difference between this text and some nonsense like “Huardest gefburn”? Kjift – not at all! A blind text like this gives you information about the selected font, how the letters are written and an impression of the look. $a\sqrt[n]{b} = \sqrt[n]{a^n b}$. This text should contain all letters of the alphabet and it should be written in of the original language. $d\Omega = \sin\vartheta d\vartheta d\varphi$. There is no need for special content, but the length of words should match the language. Hello, here is some text without a meaning. This text should show what a printed text will look like at this place. $\sin^2(\alpha) + \cos^2(\beta) = 1$. If you read this text, you will get no information $E = mc^2$. Really? Is there no information? Is there a difference between this text and some nonsense like “Huardest gefburn”? Kjift – not at all! A blind text like this gives you information about the selected font, how the letters are written and an impression of the look. $\sqrt[n]{a} \cdot \sqrt[n]{b} = \sqrt[n]{ab}$. This text should contain all letters of the alphabet and it should be written in of the original language. $\frac{\sqrt[n]{a}}{\sqrt[n]{b}} = \sqrt[n]{\frac{a}{b}}$. There is no need for special content, but the length of words should match the language. $a\sqrt[n]{b} = \sqrt[n]{a^n b}$.

1.1 A Section

Hello, here is some text without a meaning. $d\Omega = \sin\vartheta d\vartheta d\varphi$. This text should show what a printed text will look like at this place. If you read this text, you will get no information. Really? Is there no information? Is there a difference between this text and some nonsense like “Huardest gefburn”? Kjift – not at all! A blind text like this gives you information about the selected font, how the letters are written and an impression of the look. $\sin^2(\alpha) + \cos^2(\beta) = 1$. This text should contain all letters of the alphabet and it should be written in of the original language $E = mc^2$. There is no need for special content, but the length of words should match the language. $\sqrt[n]{a} \cdot \sqrt[n]{b} = \sqrt[n]{ab}$. Hello, here is some text without a meaning. $\frac{\sqrt[n]{a}}{\sqrt[n]{b}} = \sqrt[n]{\frac{a}{b}}$. This text should show what a printed text will look like at this place. $a\sqrt[n]{b} = \sqrt[n]{a^n b}$. If you read this text, you will get no information. $d\Omega = \sin\vartheta d\vartheta d\varphi$. Really? Is there no information? Is there a difference between this text and some nonsense like “Huardest gefburn”? Kjift – not at all! A blind text like this gives you information about the selected font, how the letters are written and an impression of the look. This text should contain all letters of the alphabet and it should be written in of the original language. There is no need for special content, but the length of words should match the language.

$$\sin^2(\alpha) + \cos^2(\beta) = 1.$$

- First item in a list
- Second item in a list
- Third item in a list
- Fourth item in a list
- Fifth item in a list

1. First item in a list
2. Second item in a list
3. Third item in a list
4. Fourth item in a list
5. Fifth item in a list

First item in a list

Second item in a list

Third item in a list

Fourth item in a list

Fifth item in a list

1.2 Cool Pictures

Figure 1.1: Itsuki heard some shocking information

Hello, here is some text without a meaning $E = mc^2$. This text should show what a printed text will look like at this place. $\sqrt[n]{a} \cdot \sqrt[n]{b} = \sqrt[n]{ab}$. If you read this text, you will get no information. $\frac{\sqrt[n]{a}}{\sqrt[n]{b}} = \sqrt[n]{\frac{a}{b}}$. Really? Is there no information? Is there a difference between this text and some nonsense like “Huardest gefburn”? Kjift – not at all! A blind text like this gives you information about the selected font, how the letters are written and an impression of the look. $a\sqrt[n]{b} = \sqrt[n]{a^n b}$. This text should contain all letters of the

alphabet and it should be written in of the original language. $d\Omega = \sin\vartheta d\vartheta d\varphi$. There is no need for special content, but the length of words should match the language. Hello, here is some text without a meaning. This text should show what a printed text will look like at this place. $\sin^2(\alpha) + \cos^2(\beta) = 1$. If you read this text, you will get no information $E = mc^2$. Really? Is there no information? Is there a difference between this text and some nonsense like “Huardest gefburn”? Kjift – not at all! A blind text like this gives you information about the selected font, how the letters are written and an impression of the look. $\sqrt[n]{a} \cdot \sqrt[n]{b} = \sqrt[n]{ab}$. This text should contain all letters of the alphabet and it should be written in of the original language. $\frac{\sqrt[n]{a}}{\sqrt[n]{b}} = \sqrt[n]{\frac{a}{b}}$. There is no need for special content, but the length of words should match the language. $a\sqrt[n]{b} = \sqrt[n]{a^n b}$.

I Material

- Fichte ?

II Deskplate

- amount: 1
- measurements: 2500 x 1250 x 27 [mm]

III Reinforcement

- amount: 1
- Length: 2500 mm
- Width: ≤ 90 mm
- Height: $160 > H > 140$ [mm]
- width (W) is determined by two Ikea Alex but proper legs might give more flexibility later on

Chapter 2

Chapter 3



Chapter 4



Chapter 5



Chapter 6

Creating Realistic Rendering Effects

6.1 Understanding graphics shaders

1. OpenGL shading language (GLSL) provides the ability to develop graphics shaders
→ blocks of graphic software instructions to calculate more realistic rendering effects, rather than fixed function states.
2. steps to designing shaders and applying them to a scene
 - write your own shaders ("like C programs"). They are treated as a set of strings passed to the hardware so create them on the fly or read them as text files.
 - specify no more than a vertex shader, a geometry shader and a fragment shader to be processed in the OpenGL pipeline. Each stage has only one `main()` function.
 - will totally replace fixed functionalities such as fog, lighting and texture mapping, which have to be re-implemented in your shader source code.
 - Shaders require OpenGL API to compile and execute them.
 - Vertex shader can apply transformations to each vertex
 - Fragment shaders calculate the color of individual pixels coming from the rasterizer;
 - Geometry shaders re-generate geometries from existing vertices and primitive data

6.1.1 osg::Shader

- define shader object containing source code strings.
- `setShaderSource()` specifies the src code text from a `std::string` variable
- `loadShaderSourceFromFile()` reads a source file from drive.
- construct shader object from existing string like this:

```
osg::ref_ptr<osg::Shader> vertShader =
    new osg::Shader( osg::Shader::VERTEX, vertText );
```

- input param `OSG::Shader::VERTEX` represents the vertex shader. Use `GEOMETRY` or `FRAGMENT` enums instead to specify geometry- or fragment shader.

```
osg::ref_ptr<osg::Shader> fragShader =
    new osg::Shader( osg::Shader::FRAGMENT, fragText );

osg::ref_ptr<osg::Shader> geomShader =
    new osg::Shader( osg::Shader::GEOMETRY );

geomShader -> loadShaderSourceFromFile( "source.geom" );
```

→source.geom contains geometry shader.

- `osgDB::readShaderFile()` may be even better
 - automatically checks shader types (via extensions: `.vert`, `.frag`, `.geom`)
 - returns `osg::Shader` instance of correct type and data:

```
osg::Shader* fragShader =
    osgDB::readShaderFile( "source.frag" );
```

→shaders are set and ready to be use

→use `osg::Program` calss and `addShader()` method to include include shaders and set GLSL rendering attribute and modes to a state set.

- most other fixed-function states willbecome incalid after the shaders make effects, including lights, materials, fog, texture mapping, texture coordinate generation and texture environment.
- following code adds all above shaders to an `osg::Program` object and attaches it to the state set of existing node:

```

osg::ref_ptr<osg::Program> program =
    new osg::Program;
program -> addShader( vertShader.get() );
program -> addShader( fragShader.get() );
program -> addShader( geomShader.get() );
node -> getOrCreateStateSet() -> setAttributeAndModes(
    program.get() );

```

6.2 Using uniforms

- three types of inputs and outputs in a typical shader:
 - uniforms
 - vertex attributes
 - varyings
- Uniforms and Vertex Attributes are read-only during the sahder's execution, but can be set by host OpenGL or OSG apps.
 - They are actually global GLSL variables used for interactions between shaders and user applications.
- Varyings are used for passing data from one shader to the next one
 - tehy are invisible to external programs
- OSG uses osg::Uniform class

6.2.1 osg::Uniform class

- used to define a SLGL uniform cariable
- constructor has a name and initial value param, which should match the definition in the shader souce code, e.g:

```

float length = 1.0f;
osg::ref_ptr<osg::Uniform> uniform =
    new osg::Uniform( "length", length );

```

- add uniform object to state set, which has attached osg::Program object via addUniform():

```

stateset -> addUniform( uniform.get() );

```


There should be a variable defined in one of the shader sources:

```
uniform float length;
```

Otherwise, uniform variable will not be available in either OSG programs or shaders.

- Uniforms can be any basic type, or any aggregation of types, such as Boolean, float, integer, 2D/3D/4D vector, matrix and various texture samplers.
- `osg::Uniform` class accepts all basic types with constructor and `set()` method.
→ additionally, `osg::Matrix2` and `osg::Matrix3`
- to bind texture sampler (used in shaders to represent a particular texture) you specify the texture mapping unit by using an unsigned int:

```
osg::ref_ptr<osg::Uniform> uniform =  
    new osg::Uniform( "texture", 0 );
```

- there must already be an `osg::Texture` object at unit 0, as well as a sampler uniform in the shader source:

```
uniform sampler2D texture;
```

→ assume that it's a 2D texture that will be used to change the shader's executing behavior.

6.2.2 Time for Action page 154

6.2.3 What just happened?

basic algorithm for caroon shading:

- if there's a normal that is close to the light direction, the brightest tone
→ `color1` is used.

- as the angle between light direction and surface normal is increasing
→darker tones will be used (color2, color3, color4)
→provides an intensity value for selecting tones.
→all four tones are declares as 4D vectors in FRAGMENT SHADER and passed to osg::Uniform objects as osg::Vec4 variables in the user app.

6.3 Working with the geometry shader

- geometry shader is included into the OpenGL 3.2 core
→in lower versions it is udes as an extension (GL_EXT_ geometry_shader4) which should be declared in the shader sourve code.
- geometry shader has new sdjacency primitives
→can be used as arguments of osg::PrimitiveSet derived classes.
→also requires setting up params in order to maipulate the shader operations:

1. GL_GEOMETRY_VERTICES_OUT_EXT: numns of vertices that the shader will emit
2. GL_GEOMETRY_INPUT_TYPE_EXT: the primitive type to be sent to the shader
3. GL_GEOMETRY_OUTPUT_TYPE_EXT: primitive type to be emitted from the shader

→osg::Program class's setPatameter() sets values for these params

→100 vertices will be emitted from the shader to the primitive assembly processor in the rendering pipeline:

```
program -> setParameter( GL_GEOMETRY_VERTICES_OUT_EXT, 100 );
```

6.3.1 Time for action - Generating a Bezier curve

P158

6.3.2 What just happened?

- geometry shader defines a new primitive type GL_LINE_STRIP_ADJACENCY_EXT which means a line strip with adjacency
→first and last vertices provide adjacency information bur aren't visible as line segments.

→thus we can use these two extra vertices as the endpoints of a Bezier curve and the others as control points

→that is actually what we read from the GLSL variable `gl_Position[0]` to `gl_PositionIn[3]`.

- Cubic Bezier curve equation:

$$P(t) = (1-t)^3 * P0 + 3*t*(1-t)^2 * (1-t) * P2 + t^3 * P3 \text{ with } 0 \leq t \leq 1$$

See summary.

Chapter 7

Viewing the World

Focus:

- understanding the coordinate system defined in OpenGL
- alternating the view point and orientation, projection frustum, and final viewport
- changing and controlling the rendering order if there exists more than one camera
- how to create single and composite viewers
- how to manage global display settings and generate easy-to-use stereo visualization effects
- how to apply the rendered scene as a texture object - so called rendering to textures (RTT)

7.0.1 From world to screen

this subsection will be shorter, since a version of this is already in my personal notebook.

modelmatrix

used to describe the specific location of an object in the world.

→ transforms object's local coord sys to world coord sys. Both coord. systems are right-handed.

view matrix

→transforms entire world into view space. suppose we have a camera placed at a certain position in the world; the inverse of the camera's transformation matrix is actually used as the view matrix.

In the right-handed view coord sy, OpenGL defines that the camera is always located at the origin (0, 0, 0), and facing towards the negative Z axis.

→Hence, we can represent the world on our camera's screen.

Note:

There is no separate model matrix or view matrix in Open GL.

→however, it defines a model-view matrix to transform from the object's local space to view space, which is a combination of both matrices.

→to transform vertex V in local space to V_e in view space, we have:

$$V_e = V * modelViewMatrix$$

projection matrix

we have to:

- determine how 3D objects are projected onto the screen (perspective or orthogonal)
- calculate the frustum.

→Projection matrix is used to specify the frustum in the world coordinate system with six clipping planes: left, right, bottom, top, near and far planes.

→OpenGL function: `gluPerspective()`, determines a field of view with camera lens params.

- resulting coord sys is called: Normalized Device Coordinate System
 - it ranges from -1 to +1 in each of the axes.
 - is changed to left-handed now.
- as a final step:
 - project all result data onto viewport. (the window)
 - define the window rectangle in which the final image is mapped
 - As well as Z Value of the window coordinates.
- Now the 3D scene is rendered to a rectangular area on your 2D screen.

MVPW matrix

Finally, the screen coord Vs can represent the local vertex V in the 3D world by using the so called MVPW matrix:

$$Vs = V * modelViewMatrix * projectionMatrix * windowMatrix$$

The Vs is still a vector that represents a 2D pixel location with a depth value.

By reversing this mapping process, we can get a line in the 3D space from a 2D screen point (Xs, Ys)

→that's because th 2D point can actually be treated as two points: one on the near clipping plane ($Zs = 0$) and the other on the far plane ($Zs = 1$).

The inverse of the MVPW matrix is used to obtain the result of the "unproject" work:

$$V0 = (Xs, Ys, 0) * invMVPW$$

$$V1 = (Xs, Ys, 1) * invMVPW$$

7.1 The Camera class

- it's popular to use `glTranslate()` and `glRotate()`
→moves the scene
- it's popular to use `gluLookAt()`
→moves the camera
- though they are all replaceable by `glMultMatrix()`
→in fact, these functions do the same thing: calculate the model-view matrix for transforming data from world space to view space.
- similarly, OSG had `osg::Transform` class
→adds or sets its own matrix to the current model-view matrix when placed in the sg
- BUT: we always intend to operate on model matrix by using the
→`osg::MatrixTransform` and `osg::PositionAttitudeTransform` subclasses
→we handle the view matrix with the `osg::Camera` subclass.
- `osg::Camera` class is one of the most important classes in the core OSG libraries.
→can be used as Group node

- but it is far more than a common node
 - main functionalities in four categories:
1. `osg::Camera` class handles the view matrix projection matrix and viewport
 - affects all its children and project them onto the screen
 - Related methods:
 - public: `setViewMatrix()` and `setViewMatrixAsLookAt()` methods set the view matrix by using the `osg::Matrix` variable or classic eye/center/up variables.
 - public `setProjectionMatrix()` method accepts an `osg::Matrix` parameter in order to specify the projection matrix
 - other convenient methods:
 - `setProjectionMatrixAsFrustum()`
 - `setProjectionMatrixAsOrtho()`
 - `setProjectionMatrixAsOrtho2D()`
 - `setProjectionMatrixAsPerspective`
 are used to set a perspective or orthographic projection matrix with different frustum parameters.
 they work just like the OpenGL projection functions (... , see page 165)
 - public `setViewport()` method defines a rectangular window area with an `osg::Viewport` object.

set view and projection matrix of a camera node, set its viewport to $(x, y) - (x + w, y + h)$:

```
camera -> setViewMatrix( viewMatrix );
camera -> setProjectionMatrix( projectionMatrix );
camera -> setViewport( new osg::Viewport( x, y, w, h ) );
```

Obtain current view and projection matrices and viewport of the `osg::Camera` object by using the corresponding `get*()` methods at any time, e.g.:

```
osg::Matrix viewMatrix = camera -> getViewMatrix();
```

get position and orientation of view matrix:

```
osg::Vec3 eye, center, up;
camera -> getViewMatrixAsLookAt( eye, center, up );
```

2. `osg::Camera` encapsulates the OpenGL functions, such as `glClear()`, `glClearColor()`, and `glClearDepth()`, and clears the frame buffers and presets their values when redrawing the scene to the window

in every frame.

Primary methods include:

- `setClearMask()` method, sets buffer to be cleared.
default:

`GL_COLOR_BUFFER_BIT | GL_DEPTH_BUFFER_BIT`

- `setClearColor()` method sets the clear color in RGBA format, by using an `osg::Vec4` variable.
 - similarly there's `setClearDepth()`, `setClearStencil()`, `setClearAccum()` (and their `get*()` methods)
3. third category includes the management of OpenGL graphics context associated with this camera.
→Chapter 9 Interacting with Outside Elements
 4. Finally, a camera can attach a texture object to internal buffer components (color buffer, depth buffer, and so on) and directly render the sub-scene graph into this texture.
→the resultant texture can then be mapped to surfaces of other scenes. This technique is named render-to-textures or texture baking
→later this chapter.

7.1.1 Rendering order of cameras

- at least one main camera node in any sg.
→created and managed by the `osg::Viewer` class
→read it with `getCamera()` method.
- It automatically adds the root node as its child node before starting the simulation.
→by default all other cameras (directly and indirectly added to root node) will share the graphics context associated with the main camera, will share the graphics context associated with the main camera + draw their own sub-scenes successively onto the same rendering window.
- `osg::Camera` class provides `setRenderOrder()` method to precisely control the rendering order of cameras.
→It has an order enum and an optional order num param.
→first enum is either `PRE_RENDER` or `POST_RENDER` (indicates general rendering order)
→second is an interger num for sorting cameras of the same type in ascending order. (default = 0)

- this will force OSG to render camer1 first, then camera2 (larger int num), then camera3:

```
camera1 -> setRenderOrder( osg::Camera::PRE_RENDER );
camera2 -> setRenderOrder( osg::Camera::PRE_RENDER, 5 );
camera3 -> setRenderOrder( osg::Camera::POST_RENDER );
```

If a camera is rendered first (**PRE_RENDER**) its rendering result in the buffers will be cleared and covered by the next camera, and the viewer may not be able to see its sub-scene. This is especially useful in the case of the render-to-textures process, because we want the sub-scene to be hidden from the screen, and update the attached texture objects before starting the main scene.

In addition, if a camera is rendered afterwards (**POST_RENDER**), it may erase the current color and depth values in the buffers.
→ avoid this by calling `setClearMask()` with fewer buffer masks. (HUD head up display)

7.1.2 Time for action creating an HUD camera

P168

7.1.3 WTF just happened?

an additional camera contains the glider model that is to be rendered as its sub-scene-graph on top of the rendering result (color buffer and depth buffer) of the main camera.

The additional camera's goal is to implement a HUD scene that overlays the main scene. It clears the depth buffer to ensure that all pixel data drawn by this camera can pass the depth test. However, the color buffer is not cleared, keeping the uncovered pixel data of the main scene on the screen. That is why we set it up like this;

```
camera -> setClearMask( GL_DEPTH_BUFFER_BIT ); // no color buffer bit
```

7.1.4 using a single viewer

OSG supports the single viewer class `osgViewer::Viewer` for holding a view on a single scene.

`setSceneData()` method

→ manages the scene graph's root node

`run()`

→starts the simulation loop (scene is rendered per frame) →the frame buffer is updated continuously by the result of every rendering cycle (-i frame)

the viewer also contains an `osg::Camera` object as the main camera.

View Matrix of the camera is controlled by the viewer's internal `osgGA::CameraManipulator` object.

User input events are received and handled by the viewer as well..

→this works via a list of `osgGA::GUIEventHandler` handlers.

The viewer can even be set up in full screen mode, in a window and onto a spherical display.

7.1.5 Digging into the simulation loop

The simulation loop defined by the `run()` method always has three types of tasks to perform:

1. specify the main camera's manipulator
2. set up associated graphics contexts
3. render frames in cycles

The manipulator can read keyboard and mouse events and accordingly adjust the main camera's view matrix to navigate the scene graph.

→set by using `setCameraManipulator()` method

→param: `osgGA::CameraManipulator` subclass

e.g.:

```
viewer.setCameraManipulator( new osgGA::TrackballManipulator );
```

→adds trackball(arc ball) manip to viewer object, (free motion behaviour)

→because camera manipulator is kept a smart pointer in the viewer, we can assign a new manip by using the `setCameraManipulator()` method any time.

see page 170 for table with manipulators

→beware, to declare and use a manip you should add the `osgGA` library as a dependence of your project

→CMake scripts

the graphics contexts of a viewer, as well as possible threads and resources, are all initialized in the `realize()` method

→automatically called before the first frame is rendered

Now the viewer enters the loop:

→each time `frame()` method is used to render a frame. It checks if the rendering process should stop and exit with the `don()` method. The process can be described with just a few lines of code:

```
while( !viewer.done() )
{
    viewer.frame();
}
```

→default rendering scheme used in the viewer class.

Frame rate is synched to the monitor's refresh rate to avoid wasting system energy, if the vsync option of the gpu is on.

OSG supports another on-demand rendering scheme:

```
viewer.setRunFrameScheme( osgViewer::Viewer::ONDEMAND );
```

now the `frame()` method will only be executed when there are scene graph mods, updating processes, or user input events, until the scheme is changed back to the default value of `CONTINUOUS`.

As an addition, the `osgViewer::Viewer` class also contains a `setRunMaxFrameRate()` method which uses a frame rate number as the param.

→can set a max frame rate

→controls viewer running to force rendering frames without lots of consumption.

7.1.6 Time for action - custom simulation loop

`run()` was used many times.

→performed update, cull and draw traversals each frame.

→see P172.

7.1.7 What on earth just happened?

this was the concept of pre- and post-frame events and assume they are executed before and after `frame()` method.

→inaccurate.

multiple threads are used to manage user updating, culling and drawing of different cameras.

→especially with multiple screens, processors, and gpu.

`frame()` method only starts a new updating/culling/drawing traversal work, but does not take care of thread synchronization. In this case, the code before and after `frame()` will be considered unstable and unsafe, because they may

conflict with other process threads when reading or writing the scene graph.
 →so the approach described here is not recommended for future development.
 →”correct” methods in next chapter.

When will the `viewer.done()` return true?

Of course, you can set the done flag via `setDone()` method of `viewer`. The OSG system will check if all present graphics contexts (for example, the rendering window) have been closed, or if the Esc key is pressed which will also change the done flag.

`viewer.setEventSetsDone()` method can even set which key is going to carry out the duty instead of the default Esc (or set this to 0 to turn off the feature).

7.1.8 Using a composite viewer

`osgViewer::Viewer` class manages only one single view on one scene graph.

`osgViewer::CompositeViewer` class supports multiple views and multiple scenes.

This has the same methods such as `run()`, `frame()` and `done()` to manage the rendering process, but also supports adding and removing independent scene views by using the `addView()` and `removeView()` methods, and obtaining a view object at a specific index by using the `getView()` method. The view object here is defined by the `osgViewer::View` class.

`osgViewer::View` class is the super class of `osgViewer::Viewer`

→it accepts setting a root node as the scene data, and adding a camera manipulator and event handlers to make use of user events as well.

The main difference between `osg::Viewer::View` and `osgViewer::Viewer` is that the former cannot be used as a single viewer directly - that is, it doesn't have `run()` or `frame()` methods.

To add a created view object to the composite viewer:

```
osgViewer::CompositeViewer multiviewer;
multiviewer.addView( view );
```

7.1.9 Time for action - rendering more scenes at one time

Multi-viewers are practical in representing complex scenes, for instance, to render a wide area with a main view and an eagle eye view, or to display the front, side, top, and perspective views of the same scene. Here we will create

three separate windows, containing three different models, each of which can be independently manipulated.

7.1.10 What the fuck just happened?

it's possible to create three `osg::Camera` nodes, add different sub-scenes to them, and attach them to different graphics contexts (rendering window) in order to achieve the same result as the previous image.

→every `osgViewer::View` object has an `osg::Camera` node that can be used to manage its subscene and associated window.

→it actually works like a container.

However, the `osgViewer::View` class handles manipulator and user events, too.

→in a composite viewer, each `osgViewer::View` object holds its own manipulator and event handlers (this will be discussed in Chapter 9, Interacting with Outside Elements).

However, a set of cameras can hardly interact with user inputs separately. That is why we choose to use a composite viewer and a few view objects to represent multiple scenes in some cases.

7.1.11 Changing global display settings

OSG manages a set of global display settings that are required by cameras, viewers, and other scene elements. It uses the singleton pattern to declare a unique instance of the container of all of these settings, by using the `osg::DisplaySettings` class. We can thus obtain the display settings instance at any time in our apps:

```
osg::DisplaySettings* ds = osg::DisplaySettings::instance();
```

The `osg::DisplaySettings` instance sets up properties requested by all newly created rendering devices, mainly OpenGL graphics contexts of rendering windows. Its characteristics include:

1. `setDoubleBuffer()` method: set double or single buffering. Default is on.
2. `setDepthBuffer()` method: whether to use depth buffer or not. Default is on.
3. `setMinimumNumAlphaBits()` (and others): set bits for an OpenGL alpha buffer, a stencil buffer, accumulation buffer. Defaults are all 0.

4. `setNumMultiSamples()`: set using multisampling buffers and number of samples. default is 0.
5. enable stereo rendering and configure stereo mode and eye mapping parameters

→some of these characteristics can be separately set for different graphics contexts by sing a specific traits structure. For now we use global display settings.

7.1.12 Time for action - enabling global multisampling

P180

7.1.13 wth just happened?

multisampling allows apps to create a frame buffer with a given number of samples per pixel.

→contains color depth, stencil info.

→more video memory is required but a better rendering result will be produced.

OSG has an internal graphics context manager `osg::GraphicsContext`:

→it's subclass `osg::GraphicsWindowWin32` (look up Linux version) manages the config and creation of rendering windows under Windows.

It will apply these two attributes to the encapsulated `wglChoosePixelFormatARB()` function and enable multisampling of the entire scene.

`osg::DisplaySettings` actually works like a default value set of various display attributes. If there is no separate setting for a specific object, the default one will take effect; Otherwise the `osg::DisplaySettings` instance will not be put to use.

We are going to talk about the separate settings for creating graphics context and the `osg::GraphicsContext` class in Chapter 9

7.1.14 Stereo visualization

We have already experienced the charm of stereoscopic 3D films and photographs.

→James Cameron's Avatar

Anaglyph image is the earliest and most popular method of presenting stereo visualiation.

others: NVIDIA's quad-buffering, horizontal or vertical split, horizontal or vertical interlace, ...

Fortunately, OSG supports most of these common stereo techniques, and can immediately realize one of them in the viewer with just a few commands:

```
osg::DisplaySettings::instance() -> setStereoMode( mode );
osg::DisplaySettings::instance() -> setStereo( true );
```

The method `setStereoMode()` selects a stereo mode from a set of enumerations, and the `setStereo()` method enables or disables it. Available stereo modes in OSG are: `ANAGLYPHIC`, `QUAD_BUFFER (NVIDIA)`, `HORIZONTAL_SPLIT`, `VERTICAL_SPLIT` (DLP projector).

You may also use `LEFT_EYE` or `RIGHT_EYE` to indicate that the screen is used for left-eye or right-eye views.

for more stereo params, such as the eye separation, have a look at the API documentation.

7.1.15 Time for action - rendering naglyph stereo scenes

P183

7.1.16 wtf just happened?

in the `ANAGLYPHIC` mode, the final rendering result is always made up of two color layers, with a small offset to produce a depth effect. Each eye of the glasses will see a slightly different picture, and their composition produces a stereograph image, which will be fused by our brain into a three dimensional scene.

OSG supports the anaglyphic stereo mode with a two-pass rendering scheme. The first pass renders the left eye image with a red channel color mask, and the second pass renders the right eye image with a cyan channel. the color mask is defined by the rendering attribute `osg::ColorMask`. It can be easily applied to state sets of nodes and drawables by using:

```
osg::ref_ptr<osg::ColorMask> colorMask = new osg::ColorMask;
colorMask -> setmask( true, true, true, true );
stateset -> setAttribute( colorMask.get() );
```

stereo mode often causes the scene graph to be rendered multiple times, which slows down the frame rate as a side effect.

7.1.17 Rendering to textures

the render to textures technique allows developers to create textures based on a sub-scene's appearance in the rendered scene. These textures are then

”baked” into objects of coming `osg` via texture mapping. They can be used to create nice specular effects on the fly, or can be stored for subsequent deferred shading, multi-pass rendering, and other advanced rendering algorithms. To implement texture baking dynamically, there are generally three steps to follow:

1. Create the texture for rendering.
2. Render the scene to the texture.
3. Use the texture as you want.

We have to create an empty texture object before putting it into use. `OSG` can create an empty `osg::Texture` object by specifying its size. The `setTextureSize()` method defines the width and height of a 2D texture, and an additional depth parameter of a 3D texture.

The key to rendering a scene graph to the newly created texture is the `attach()` method of the `osg::Camera` class. This accepts the texture object as an argument, as well as a buffer component parameter, which indicates which part of the frame buffer will be rendered to the texture. For example, to attach the color buffer of a camera’s sub-scene to the texture, we use:

```
camera -> attach ( osg::Camera::COLOR_BUFFER, texture.get() );
```

Other usable buffer components include the `DEPTH_BUFFER`, `STENCIL_BUFFER` and `COLOR_BUFFER0` to `COLOR_BUFFER15` (multiple render target outputs, depending on the gpu).

Continue setting suitable view and projection matrices of this camera, and a viewport to meet the texture size, and set the texture as an attribute of nodes or drawables. The texture will be updated with the camera’s rendering result in every frame, dynamically carrying with the alteration of the view matrix and the projection matrix.

Be aware that the main camera of a viewer is not suitable for attaching a texture. Otherwise there will be no outputs to the actual window, which will make the screen pitch-dark. Of course, you may ignore this if you are doing off-screen rendering and don’t care of any visual effects.

7.1.18 Frame buffer, pixel buffer, and FBO

A concern is how to get the rendered frame buffer image into the texture object. A direct approach is to use the `glReadpixels()` method to return pixel data from the frame buffer, and apply the result to a `glTexImage*()` method. →easy to conceptualize and use, but will always copy data to the texture

object, which is extremely slow.

The `glCopyTexSubImage()` method would be better in terms of improving the efficiency. However, we can still optimize the process. Rendering the scene directly to a target other than the frame buffer is a good idea. There are mainly two solutions for this:

1. the pixel buffer (pbuffer for short) extension can create an invisible rendering buffer with a pixel format descriptor, which is equivalent to a window. It should be destroyed after being used, as is done for the rendering window.
2. The frame buffer object (FBO for short), which is sometimes better than pixel buffer in saving the storage space, can add application-created frame buffers and redirect the rendering output to it. It can either output to a texture object or a renderbuffer object, which is imply a data storage object.

OSG supports making use of different render target implementations: directly copying from the frame buffer pixel buffer, or FBO. It uses the method `setRenderTargetImplementation()` of the `osg::Camera` class to select a solution from them, for example:

```
camera -> setRenderTargetImplementation( osg::Camera::FRAMEBUFFER );
```

This indicates that the rendering result of Camera will be rendered to the attached texture by using the `glCopyTexSubImage()` method internally. In fact, this is the default setting of all camera nodes.

Other major implementations include `PIXEL_BUFFER` and `FRAME_BUFFER_OBJECT`

7.1.19 TFA - drawing aircrafts on a loaded terrain

In this section , we are going to integrate what we learned before to create a slightly complex example, which identifies all texture objects in a scene graph by using the `osg::NodeVisitor` utility, replaces them with a newly created shared texture, and binds the new texture to a render-to-textures camera. the texture is expected to represent more than a static image, so a customized simulation loop will be used to animate the sub-scene graph before calling the `frame()` method.

7.1.20 Summary

est